

# Regression Analysis of the Runge Function: OLS, Ridge, Lasso, and Resampling Techniques

Christopher A. Trotter

*Department of Mathematics, University of Oslo*

Oslo, Norway

Email: chrisatrotter@gmail.com

**Abstract**—This project investigates regression methods for modeling the Runge function,  $f(x) = \frac{1}{1+25x^2}$ , with Gaussian noise. We implement Ordinary Least Squares (OLS), Ridge, and Lasso regression using polynomials up to degree 15 for sample sizes  $n = 100, 500, 1000$ . Optimization employs gradient descent (GD) variants (vanilla, momentum, AdaGrad, RMSprop, ADAM) and stochastic gradient descent (SGD). Model performance is evaluated using bootstrap resampling and k-fold cross-validation, with bias-variance trade-off analysis. OLS overfits above degree 10 due to Runge's phenomenon, while Ridge achieves the lowest test MSE (0.09 at  $d = 12$ ,  $\lambda = 0.01$ ). Lasso promotes sparsity but yields higher MSE due to feature correlations. ADAM-SGD converges fastest, and larger  $n$  reduces variance. These results highlight the effectiveness of regularization and resampling for improving generalization on non-linear problems [1].

**Index Terms**—Regression Analysis, Runge Function, OLS, Ridge, Lasso, Gradient Descent, Stochastic Gradient Descent, Bootstrap, Cross-Validation, Bias-Variance Trade-off

## I. INTRODUCTION

Regression analysis is a cornerstone of statistical modeling, widely used for function approximation [2]. The Runge function,  $f(x) = \frac{1}{1+25x^2}$ , poses challenges due to Runge's phenomenon, where high-degree polynomial fits oscillate near interval boundaries [1]. This project implements Ordinary Least Squares (OLS) [3], Ridge [4], and Lasso regression [5] to model the Runge function with added Gaussian noise. We use gradient descent (GD) variants (vanilla, momentum, AdaGrad, RMSprop, ADAM) [6] and stochastic gradient descent (SGD) [7] for optimization. Model robustness is assessed via bootstrap resampling and k-fold cross-validation, including bias-variance decomposition [8]. We critically evaluate numerical stability, the impact of data scaling, and the suitability of linear models for this non-linear problem.

## II. THEORY

### A. Ordinary Least Squares (OLS)

OLS minimizes the sum of squared errors,  $\min_{\theta} \|y - X\theta\|_2^2$ . The analytical solution is  $\theta = (X^T X)^{-1} X^T y$ , computed using `np.linalg.pinv` for stability [3].

### B. Ridge Regression

Ridge adds an L2 penalty,  $\min_{\theta} \|y - X\theta\|_2^2 + \lambda \|\theta\|_2^2$ , improving stability for multicollinear data. The solution is  $\theta = (X^T X + \lambda I)^{-1} X^T y$  [4].

### C. Lasso Regression

Lasso uses an L1 penalty,  $\min_{\theta} \|y - X\theta\|_2^2 + \lambda \|\theta\|_1$ , promoting sparsity. It requires numerical optimization due to non-differentiability [5], [8].

### D. Gradient Descent (GD)

GD iteratively updates parameters:  $\theta_{t+1} = \theta_t - \eta \nabla J(\theta)$ , where  $\eta$  is the learning rate. Variants include momentum, AdaGrad, RMSprop, and ADAM [6], [9].

### E. Stochastic Gradient Descent (SGD)

SGD uses mini-batches to approximate the gradient, reducing computational cost but introducing variance [7], [9].

### F. Bias-Variance Decomposition

The expected MSE is decomposed as:

$$\mathbb{E}[(y - \tilde{y})^2] = \mathbb{E}[(y - \mathbb{E}[\tilde{y}])^2] + \mathbb{E}[(\mathbb{E}[\tilde{y}] - \tilde{y})^2] + \sigma^2,$$

where the terms are bias, variance, and irreducible noise, respectively. We approximate  $f(x) \approx y$  for bootstrap analysis [8], [10].

### G. Resampling Methods

Bootstrap resamples with replacement ( $B = 100$ ). K-fold cross-validation ( $k = 5$ ) splits data to estimate generalization error [8].

TABLE I  
BEST TEST MSE FOR RUNGE FUNCTION ( $n = 100$ ).

Method	MSE	Degree	$\lambda$
OLS	0.12	8	-
Ridge	0.09	12	0.01
Lasso	0.15	10	0.01

### III. METHODS

#### A. Data Preprocessing

We generate  $n = 100, 500, 1000$  points for the Runge function in  $[-1, 1]$  with  $\mathcal{N}(0, 0.05^2)$  noise. The design matrix uses polynomials up to degree 15. Data are split 80/20 (train/test) with `random_state=1993`. Features are standardized using `StandardScaler`, with an option to disable scaling (`-noscale`) [11].

#### B. Regression Implementation

OLS and Ridge use analytical solutions via `np.linalg.pinv`. Lasso uses `scikit-learn.Lasso` (analytical) and GD (numerical). See Appendix A [12].

#### C. Optimization

GD uses  $\eta = 0.00001$ , with variants (vanilla, momentum, AdaGrad, RMSprop, ADAM). SGD uses batch size 32. Gradient clipping prevents divergence [9].

Resampling Bootstrap uses  $B = 100$  resamples. K-fold cross-validation ( $k = 5$ ) with `KFold` tunes  $\lambda \in [10^{-5}, 10^2]$  for Ridge and Lasso [12].

Validation Implementations were tested against a linear function ( $f(x) = 2x + 1$ ) with known coefficients, achieving  $MSE < 10^{-4}$ , ensuring correctness [12].

### IV. RESULTS

Results are generated by running scripts in `code/src` (see `README.md` at <https://github.com/chrisatrotter/FYS-STK3155>).




Fig. 6. Ridge regression with SGD for Runge function ( $n = 1000$ ).




Fig. 1. OLS regression MSE and  $R^2$  vs. polynomial degree for Runge function ( $n = 100$ ).




Fig. 2. Ridge regression MSE and  $R^2$  vs. polynomial degree for Runge function ( $n = 100$ ,  $\lambda = 0.01$ ).




Fig. 7. Cross-validation MSE for OLS, Ridge, and Lasso ( $n = 100$ ).




Fig. 3. Lasso regression MSE and  $R^2$  vs. polynomial degree for Runge function ( $n = 100$ ,  $\lambda = 0.01$ ).




Fig. 4. Ridge regression with SGD for Runge function ( $n = 100$ ).




Fig. 5. Ridge regression with SGD for Runge function ( $n = 500$ ).




Fig. 9. Cross-validation MSE for OLS, Ridge, and Lasso ( $n = 1000$ ).




Fig. 10. Bootstrap bias, variance, and MSE for OLS ( $n = 100$ ).




Fig. 8. Cross-validation MSE for OLS, Ridge, and Lasso ( $n = 500$ ).




Fig. 11. Bootstrap bias, variance, and MSE for OLS ( $n = 500$ ).




Fig. 12. Bootstrap bias, variance, and MSE for OLS ( $n = 1000$ ).

### A. OLS, Ridge, and Lasso

Table I shows Ridge achieves the lowest test MSE (0.09 at  $d = 12$ ,  $\lambda = 0.01$ ), followed by OLS (0.12 at  $d = 8$ ) and Lasso (0.15 at  $d = 10$ ). Fig. 1 shows OLS overfitting above  $d = 10$ , with coefficients growing exponentially (up to  $10^5$ ). Fig. 2 demonstrates Ridge's stability, with smaller coefficients. Lasso (Fig. 3) achieves 10% sparsity at  $d = 10$ , limited by feature correlations.

### B. Gradient Descent and SGD

Table II shows ADAM converges fastest (5000 epochs, MSE 0.08) for Ridge at  $d = 5$ ,  $n = 5000$ . SGD with ADAM is 3x faster but has higher variance (Figs. 4–6). Batch size 64 reduces variance by 15% compared to 32.

TABLE II  
CONVERGENCE EPOCHS FOR GRADIENT DESCENT METHODS ( $d = 5$ ,  
 $n = 5000$ , RIDGE).

Method	Epochs
Vanilla GD	8000
Momentum	6500
AdaGrad	6000
RMSprop	5500
ADAM	5000

### C. Bias-Variance and Cross-Validation

Bootstrap (Figs. 10–12) shows low bias for  $d \geq 5$ , with variance spiking above  $d = 10$ . Train/test MSE (Fig. ??) confirms overfitting. Cross-validation (Figs. 7–9) yields MSE within 5% of bootstrap, with Ridge optimal at  $\lambda = 0.01$ .

## V. DISCUSSION

### A. Preprocessing Impacts

Standardization reduces the condition number from  $10^{10}$  to  $10^4$ , preventing `np.linalg.LinAlgError` for  $d > 10$ . Ridge further stabilizes by adding  $\lambda I$  to  $X^T X$ . An 80/20 split balances training and evaluation; 70/30 yields similar MSE [8], [11].

### B. Runge Function Results

OLS overfits above  $d = 10$  due to Runge's phenomenon, with coefficient magnitudes reaching  $10^5$  (Fig. 1) [1]. Ridge reduces MSE by 25% at  $d = 12$ ,  $\lambda = 0.01$ , by shrinking coefficients (Fig. 2). Testing  $\lambda \in [10^{-5}, 10^2]$  shows  $\lambda = 0.01$  optimal; larger  $\lambda$  over-regularizes. Lasso's sparsity is limited (10% zero coefficients at  $d = 10$ ) due to correlated polynomial features (Fig. 3) [5]. Larger  $n$  reduces variance but not bias [8].

### C. Optimization Methods

Vanilla GD is slow (8000 epochs) due to small  $\eta = 0.00001$ , chosen to avoid divergence (larger  $\eta = 0.001$  failed). Momentum and ADAM reduce epochs by 19–38% (Table II). SGD with batch size 64 reduces variance by 15% vs. 32 (Figs. 4–6). Scikit-learn's Lasso matches GD results [7], [9].

### D. Bias-Variance and Cross-Validation

Bootstrap confirms low bias ( $d \geq 5$ ) and high variance ( $d > 10$ ), aligning with train/test MSE (Fig. ??) [10]. Cross-validation selects  $\lambda = 0.01$  for Ridge, with MSE within 5% of bootstrap (Figs. 7–9). Larger  $n$  reduces variance, but non-linearity limits performance [8].

### E. Limitations

Runge's non-linearity causes oscillations, limiting linear models [1]. Lasso's sparsity is reduced by feature correlations [5]. Future work could explore kernel methods or neural networks [9].

## VI. CONCLUSION

Ridge regression best mitigates overfitting, achieving MSE 0.09 at  $d = 12$ ,  $\lambda = 0.01$  [4]. OLS and Lasso underperform due to overfitting and limited sparsity, respectively [3], [5]. ADAM-SGD converges fastest, and larger  $n$  reduces variance [7]. Linear models are limited by Runge's non-linearity; non-linear models are recommended [1]. Code: <https://github.com/chrisatrotter/FYS-STK3155>.

## VII. ACKNOWLEDGMENTS

I used AI to structure the code into modular scripts (`part_a_ols.py`, etc.) and generate the initial LaTeX report based on the IEEE template. AI assisted in organizing figures and data, ensuring compliance with assignment requirements [13].

## REFERENCES

24

- [1] C. Runge, "Über empirische funktionen und die interpolation zwischen äquidistanten ordinaten," *Zeitschrift für Mathematik und Physik*, vol. 46, pp. 224–243, 1901.
- [2] X. Yan and X. Su, *Linear Regression Analysis: Theory and Computing*, World Scientific Publishing, 2009.
- [3] J. M. Wooldridge, *The Simple Regression Model*, 4th ed. Mason, OH: Cengage Learning, 2008, pp. 22–67.
- [4] D. E. Hilt, D. W. Seegrist, U. S. F. Service, and P. Northeastern Forest Experiment Station (Radnor, "Ridge: A computer program for calculating ridge regression estimates," Dept. of Agriculture, Forest Service, Northeastern Forest Experiment Station, Technical Report 236, 1977, caption title. [Online]. Available: <https://www.biodiversitylibrary.org/item/137258>
- [5] F. Santosa and W. W. Symes, "Linear inversion of band-limited reflection seismograms," *SIAM Journal on Scientific and Statistical Computing*, vol. 7, no. 4, pp. 1307–1330, 1986.
- [6] C. Lemaréchal, "Cauchy and the gradient method," in *Optimization Stories*, 1st ed., ser. Documenta Mathematica Series, M. Grötschel, Ed. EMS Press, 2012, vol. 6, pp. 251–254, archived from the original (PDF) on 2018-12-29. Retrieved 2020-01-26.
- [7] L. Bottou, "Online algorithms and stochastic approximations," in *Online Learning and Neural Networks*. Cambridge University Press, 1998.
- [8] T. Hastie, R. Tibshirani, and J. Friedman, *The Elements of Statistical Learning*. New York: Springer, 2009.
- [9] I. Goodfellow, Y. Bengio, and A. Courville, *Deep Learning*. MIT Press, 2016, <http://www.deeplearningbook.org>.
- [10] W. N. v. Wieringen, "Lecture notes: Statistics for high-dimensional data," <https://www.few.vu.nl/~wvanwie>, 2015, accessed: 2025-10-09.
- [11] scikit-learn developers, "Preprocessing data," <https://scikit-learn.org/stable/modules/preprocessing.html>, 2025, accessed: 2025-10-09.
- [12] C. A. Trotter, "Fys-stk3155," <https://github.com/chrisatrotter/FYS-STK3155>, 2025, GitHub repository.
- [13] M. Shell, "How to use the IEEEtran class," [https://bigdataieee.org/BigData2020/files/IEEEtran\\_HOWTO.pdf](https://bigdataieee.org/BigData2020/files/IEEEtran_HOWTO.pdf), IEEE / IEEEtran Project, Technical Report / Manual, 2002, accessed: October 9, 2025.

## APPENDIX

```

1 import numpy as np
2 from sklearn.metrics import mean_squared_error,
3     r2_score
4 from sklearn.linear_model import Lasso
5
6 class RegressionModel:
7     @staticmethod
8         def ols_fit(X: np.ndarray, y: np.ndarray) -> np.
9             ndarray:
10             """Ordinary Least Squares regression."""
11             try:
12                 return np.linalg.pinv(X.T @ X) @ X.T @ y
13             except np.linalg.LinAlgError:
14                 print("Matrix inversion failed; check
15                     scaling or degree.")
16                 return np.zeros(X.shape[1])
17
18             @staticmethod
19                 def ridge_fit(X: np.ndarray, y: np.ndarray,
20                     lambda_val: float) -> np.ndarray:
21                     """Ridge regression."""
22                     XTX = X.T @ X
23                     return np.linalg.pinv(XTX + lambda_val * np.
24                         identity(XTX.shape[0])) @ X.T @ y
25
26             @staticmethod
27                 def lasso_fit(X: np.ndarray, y: np.ndarray,
28                     lambda_val: float) -> np.ndarray:
29                     """Lasso regression."""
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79

```

```

clf = Lasso(alpha=lambda_val, fit_intercept=
False, max_iter=int(1e5), tol=1e-1)
clf.fit(X, y)
return clf.coef_

@staticmethod
def predict(X: np.ndarray, theta: np.ndarray) ->
np.ndarray:
    """Predict using model coefficients."""
    return X @ theta

@staticmethod
def compute_metrics(y_true: np.ndarray, y_pred:
np.ndarray) -> tuple:
    """Compute MSE and R2 scores."""
    mse = mean_squared_error(y_true, y_pred)
    r2 = r2_score(y_true, y_pred)
    return mse, r2

@staticmethod
def gd_fit(X: np.ndarray, y: np.ndarray, eta:
float = 0.00001, max_iter: int = 10000,
tol: float = 1e-6, lambda_val: float =
0, method: str = 'vanilla',
regression_type: str = 'ridge') ->
tuple:
    """Gradient descent with various
optimization methods for OLS, Ridge, or Lasso."""
    theta = np.zeros(X.shape[1])
    v, m, s = np.zeros_like(theta), np.
zeros_like(theta), np.zeros_like(theta)
    t, beta1, beta2, gamma, epsilon = 1, 0.9,
0.999, 0.9, 1e-8

    for i in range(max_iter):
        grad = -2/X.shape[0] * X.T @ (y - X @
theta)
        if regression_type == 'ridge':
            grad += 2 * lambda_val * theta
        elif regression_type == 'lasso':
            grad += lambda_val * np.sign(theta)
        grad = np.clip(grad, -1e2, 1e2)

        if method == 'vanilla':
            theta_new = theta - eta * grad
        elif method == 'momentum':
            v = gamma * v - eta * grad
            theta_new = theta + v
        elif method == 'adagrad':
            s += grad**2
            theta_new = theta - eta * grad / (np
.sqrt(s) + epsilon)
        elif method == 'rmsprop':
            s = beta2 * s + (1 - beta2) * grad
            **2
            theta_new = theta - eta * grad / (np
.sqrt(s) + epsilon)
        elif method == 'adam':
            m = beta1 * m + (1 - beta1) * grad
            s = beta2 * s + (1 - beta2) * grad
            **2
            m_hat = m / (1 - beta1**2)
            s_hat = s / (1 - beta2**2)
            theta_new = theta - eta * m_hat / (
np.sqrt(s_hat) + epsilon)
            t += 1
        else:
            raise ValueError(f"Unknown method: {
method}")

        if np.linalg.norm(theta_new - theta) <
tol:
            return theta_new, i + 1

```

```

    theta = theta_new
    return theta, max_iter

@staticmethod
def sgd_fit(X: np.ndarray, y: np.ndarray, eta: float = 0.0001, max_iter: int = 10000,
            tol: float = 1e-6, lambda_val: float = 0, method: str = 'vanilla',
            regression_type: str = 'ridge',
            batch_size: int = 32) -> tuple:
    """Stochastic gradient descent for OLS,
    Ridge, or Lasso."""
    theta = np.zeros(X.shape[1])
    n = X.shape[0]
    v, m, s = np.zeros_like(theta), np.
    zeros_like(theta), np.zeros_like(theta)
    t, beta1, beta2, gamma, epsilon = 1, 0.9,
    0.999, 0.9, 1e-8

    for i in range(max_iter):
        indices = np.random.permutation(n)
        for start in range(0, n, batch_size):
            batch_indices = indices[start:start
+ batch_size]
            X_batch = X[batch_indices]
            y_batch = y[batch_indices]
            grad = -2/X_batch.shape[0] * X_batch
            .T @ (y_batch - X_batch @ theta)
            if regression_type == 'ridge':
                grad += 2 * lambda_val * theta
            elif regression_type == 'lasso':
                grad += lambda_val * np.sign(
theta)
            grad = np.clip(grad, -1e2, 1e2)

            if method == 'vanilla':
                theta_new = theta - eta * grad
            elif method == 'momentum':
                v = gamma * v - eta * grad
                theta_new = theta + v
            elif method == 'adagrad':
                s += grad**2
                theta_new = theta - eta * grad /
                (np.sqrt(s) + epsilon)
            elif method == 'rmsprop':
                s = beta2 * s + (1 - beta2) *
                grad**2
                theta_new = theta - eta * grad /
                (np.sqrt(s) + epsilon)
            elif method == 'adam':
                m = beta1 * m + (1 - beta1) *
                grad
                s = beta2 * s + (1 - beta2) *
                grad**2
                m_hat = m / (1 - beta1**t)
                s_hat = s / (1 - beta2**t)
                theta_new = theta - eta * m_hat
                / (np.sqrt(s_hat) + epsilon)
                t += 1
            else:
                raise ValueError(f"Unknown
method: {method}")
            if np.linalg.norm(theta_new - theta) < tol:
                return theta_new, i + 1
            theta = theta_new
    return theta, max_iter

```

part\_a\_ols.py, part\_b\_ridge.py,  
 part\_c\_gd.py, part\_d\_gd\_advanced.py,  
 part\_e\_lasso.py, part\_f\_sgd.py,  
 part\_g\_bootstrap.py, part\_h\_cv.py,  
 utils.py.

Challenges included managing high-degree polynomial instability and optimizing Lasso GD. Reflections: Understanding trade-offs (e.g., bias-variance, speed vs. stability) deepened my appreciation for regularization and resampling.

- 1) data/: Runge function data and predictions.
- 2) figures/: Plots in subfolders part\_a, part\_b, part\_e, part\_f, part\_g, part\_h.
- 3) code/: project1\_regression\_analysis.py,